

# **“Smoking Adjoints” fast Monte Carlo Greeks**

Mike Giles

`mike.giles@maths.ox.ac.uk`

# “Smoking Adjoints” for fast Greeks

- joint work with Paul Glasserman – appeared in Risk magazine in 2006 and now being used by some leading banks
- a particularly efficient way of implementing the pathwise sensitivity approach
- builds on lots of well-established ideas in design optimisation and optimal control theory
- ideal when wanting the sensitivity of one output to changes in many different inputs
- guaranteed to give all first order derivatives for a total cost which is less than factor 4 greater than original calculation

# Generic Problem

Suppose we have a multi-dimensional SDE with numerical approximation

$$\hat{S}_{n+1} = g_n(\hat{S}_n)$$

and we want to compute sensitivity of a European option

$$\mathbb{E} \left[ f(\hat{S}_M) \right]$$

to changes in  $S_0$  and other input parameters.

# Standard pathwise sensitivity

For the Deltas we can define

$$\hat{s}_n = \frac{\partial \hat{S}_n}{\partial S_0}$$

with  $\hat{s}_0 = I$ , and differentiate each timestep evolution to get

$$\hat{s}_{n+1} = G_n \hat{s}_n$$

We then have (under certain conditions)

$$\frac{\partial}{\partial S_0} \mathbb{E} \left[ f(\hat{S}_M) \right] = \mathbb{E} \left[ \frac{\partial f}{\partial \hat{S}_M} \hat{s}_M \right]$$

with

$$\frac{\partial f}{\partial \hat{S}_M} \hat{s}_M = \frac{\partial f}{\partial \hat{S}_M} G_{M-1} G_{M-2} \dots G_1 G_0$$

# Crucial observation

Evaluating

$$\frac{\partial f}{\partial \hat{S}_M} G_{M-1} G_{M-2} \dots G_1 G_0$$

from right to left involves a sequence of matrix-matrix products, each with  $O(D^3)$  cost where  $D$  is the dimension of the SDE.

Alternatively, evaluating the same expression from left to right involves a sequence of vector-matrix products, each with  $O(D^2)$  cost – big savings if  $D$  is large.

Important: get the same result either way, so still have usual differentiability requirements of pathwise sensitivity calc

# Adjoint formulation

Starting with

$$v_M = \left( \frac{\partial f}{\partial \hat{S}_M} \right)^T$$

the adjoint iteration is given by

$$v_n = G_n^T v_{n+1}$$

and we finish with

$$\frac{\partial}{\partial S_0} \mathbb{E} \left[ f(\hat{S}_M) \right] = \mathbb{E} \left[ v_0^T \right]$$

# Adjoint formulation

Note: we have to first do the path calculation, store everything needed for the  $G_n$ , then do the adjoint calculation of the sensitivity.

The storage requirements for a single path are minimal – the storage is then reused for the next path.

However, in PDE applications these storage issues can become significant.

# Standard pathwise sensitivity

For the Vegas we can define

$$\hat{s}_n = \frac{\partial \hat{S}_n}{\partial \sigma}$$

with  $\hat{s}_0 = 0$ , and differentiate each timestep evolution to get

$$\hat{s}_{n+1} = G_n \hat{s}_n + b_n, \quad b_n \equiv \frac{\partial g_n}{\partial \sigma}$$

We then have

$$\frac{\partial f}{\partial \hat{S}_M} \hat{s}_M = \sum_{n=0}^{M-1} \frac{\partial f}{\partial \hat{S}_M} G_{M-1} G_{M-2} \dots G_{n+2} G_{n+1} b_n$$



# Adjoint formulation

This can be re-expressed as

$$\frac{\partial f}{\partial \hat{S}_M} \hat{s}_M = \sum_{n=0}^{M-1} v_{n+1}^T b_n$$

where the adjoint variables  $v_n$  are as defined before.

Hence we finish with

$$\frac{\partial}{\partial \sigma} \mathbb{E} \left[ f(\hat{S}_M) \right] = \mathbb{E} \left[ \sum_{n=0}^{M-1} v_{n+1}^T b_n \right]$$

# Automatic Differentiation

The explanation above gives the essential ideas, but doesn't explain the guarantee that all first-order derivatives of a single output can be computed at a cost no more than 4 times greater than the original computation

In addition, in real implementations you would not really store and use the matrices  $G_n$

This brings us to an area of computer science research called **automatic differentiation** (or sometimes **algorithmic differentiation**)

# Automatic Differentiation

A computer instruction creates an additional new value:

$$\mathbf{u}^n = \mathbf{f}^n(\mathbf{u}^{n-1}) \equiv \begin{pmatrix} \mathbf{u}^{n-1} \\ f_n(\mathbf{u}^{n-1}) \end{pmatrix},$$

A computer program is the composition of  $N$  such steps:

$$\mathbf{u}^N = \mathbf{f}^N \circ \mathbf{f}^{N-1} \circ \dots \circ \mathbf{f}^2 \circ \mathbf{f}^1(\mathbf{u}^0).$$

# Automatic Differentiation

In forward mode, differentiation w.r.t. one element of the input vector gives

$$\dot{\mathbf{u}}^n = D^n \dot{\mathbf{u}}^{n-1}, \quad D^n \equiv \begin{pmatrix} I^{n-1} \\ \partial f_n / \partial \mathbf{u}^{n-1} \end{pmatrix},$$

and hence

$$\dot{\mathbf{u}}^N = D^N D^{N-1} \dots D^2 D^1 \dot{\mathbf{u}}^0$$

# Automatic Differentiation

In reverse mode, we consider the sensitivity of the last element of the output vector (the final value we care about) to get

$$\begin{aligned} (\bar{\mathbf{u}}^{n-1})^T &\equiv \frac{\partial u_i^N}{\partial \mathbf{u}^{n-1}} = \frac{\partial u_i^N}{\partial \mathbf{u}^n} \frac{\partial \mathbf{u}^n}{\partial \mathbf{u}^{n-1}} = (\bar{\mathbf{u}}^n)^T D^n, \\ \implies \bar{\mathbf{u}}^{n-1} &= (D^n)^T \bar{\mathbf{u}}^n. \end{aligned}$$

and hence

$$\bar{\mathbf{u}}^0 = (D^1)^T (D^2)^T \dots (D^{N-1})^T (D^N)^T \bar{\mathbf{u}}^N.$$

Note: need to go forward through original calculation to compute/store the  $D^n$ , then go in reverse to compute  $\bar{\mathbf{u}}^n$

# Automatic Differentiation

At the level of a single instruction

$$c = f(a, b)$$

the forward mode is

$$\begin{pmatrix} \dot{a} \\ \dot{b} \\ \dot{c} \end{pmatrix}^n = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ \frac{\partial f}{\partial a} & \frac{\partial f}{\partial b} \end{pmatrix} \begin{pmatrix} \dot{a} \\ \dot{b} \end{pmatrix}^{n-1}$$

and so the reverse mode is

$$\begin{pmatrix} \bar{a} \\ \bar{b} \end{pmatrix}^{n-1} = \begin{pmatrix} 1 & 0 & \frac{\partial f}{\partial a} \\ 0 & 1 & \frac{\partial f}{\partial b} \end{pmatrix} \begin{pmatrix} \bar{a} \\ \bar{b} \\ \bar{c} \end{pmatrix}^n$$

# Automatic Differentiation

This gives a prescriptive algorithm for reverse mode differentiation.

Again the reverse mode is much more efficient if we want the sensitivity of a single output to multiple inputs.

Key result is that the cost of the reverse mode is at worst a factor 4 greater than the cost of the original calculation, regardless of how many sensitivities are being computed!

The storage of the  $D^n$  is minor for SDEs – much more of a concern for PDEs. There are also extra complexities when solving implicit equations through a fixed point iteration.

# Automatic Differentiation

Manual implementation of the forward/reverse mode algorithms is possible but tedious.

Fortunately, automated tools have been developed, following one of two approaches:

- operator overloading (ADOL-C, FADBAD++)
- source code transformation (Tapenade, TAF/TAC++, ADIFOR)

My personal experience is with Tapenade for Fortran, and FADBAD++ for C++. Both are easy to use, Tapenade is as efficient as hand-coded, FADBAD++ less so.




# LIBOR Application

- testcase from “Smoking Adjoints” paper
- good real-world example involving stochastic evolution of future interest rates
- test problem performs  $N$  timesteps with a vector of  $N+40$  forward rates, and computes the  $N+40$  deltas and vegas for a portfolio of swaptions
- hand-coded adjoint for maximum efficiency – only about 50 lines of code so not too painful to do by hand


# LIBOR Application

Finite differences versus forward pathwise sensitivities:



# LIBOR Application

Hand-coded forward versus adjoint pathwise sensitivities:



# Closing words

- the need for efficient Greeks means this research has been picked up quickly by the banks
- adjoint approach gives one level of differentiation for very little cost
- extending it to second order derivatives, can get Hessian matrix for usual cost of first order derivatives using pathwise approach
- however, this again raises the limitations due to differentiability – this is something I am working on with a “vibrato” MC idea which combines the best of both pathwise and LRM sensitivity calculations